Disaster-Relief-Project-Part-1

Mahin Ganesan, Wyatt Priddy, and Taylor Tucker

2024-03-17

Introduction

In the aftermath of a devastating earthquake that rattled Haiti, displaced people are living in makeshift shelters awaiting support from aid workers. The aid workers, mainly from the United States military, are trying to reach the dispersed camps. With communication lines being down and the terrain being impassable, there are challenges in providing relief in a time- sensitive manner.

It is known that the makeshift shelters are largely constructed with blue tarps, therefore the Rochester Institute of Technology deployed airplanes to collect high resolution geo-referenced imagery. This will help us identify where displaced persons are based on the tarps.

To determine where to allocate aid, we need to use data mining against the thousands of photos taken each day which human eyes can not efficiently filter through. Determining the location in a timely manner will be paramount to rendering aid successfully and saving human life.

The primary aim of this experiment is to evaluate the efficacy of various classification algorithms in accurately and promptly identifying the presence of makeshift shelters and, by extension, the displaced persons residing within them. By harnessing the power of machine learning and image analysis, our objective is to develop a robust algorithm capable of rapidly scanning through the imagery data, pinpointing areas of interest, and facilitating timely intervention by rescue teams.

To facilitate efficient rescue efforts, we will need to determine a threshold where the number of false positives is minimal. This may not necessarily be the model that has the highest performance in accuracy but rather the highest performance with precision, or the proportion of true positives that are correctly identified by the model out of all true positives and false positive. If there are additional aid resources after rescuing all persons identified from our models, we can loosen our model specifications to classify more objects as blue tarps in an attempt to expand our search efforts.

Data

Team members have assisted our mission by classifying training data consisting of 63,241 data points for our investigation. There are 5 classifications that have been assigned to the pixel level data:

1. Blue Tarp
2. Rooftop
3. Soil
4. Various Non-Tarp
5. Vegetation

As expected when trying to find a needle in a haystack, our representation of misplaced persons (blue tarps) makes up only a small portion of the data set. Just 3.2% of the total data set, or 2022 records, are classified as `blue tarp`. After inspecting the average color of the classes, it becomes apparent that even though blue tarps are a small section of the data their distinction in color should set them apart.

Figure 1: Average Color of Class

Vegetation and soil cover over 73% of the pictures as to be expected of pictures that largely will include countryside.


Figure 2: Distribution of Classifications in Training Set

Description of Methodology

Given that our sample data set only has 3.2% of the target class, we will use stratification to account for the imbalance in our test and train split. We are looking to ensure that there is even representation between our two sets to ensure a robust analysis when training our models and subsequently predicting on the test data. We are using an 80/20 split on the training versus testing data. Since the data set is relatively large and we are using stratification to ensure an even population of blue tarps across the split, leaving 20% of the data for hold-out is appropriate.

Each model is being trained with 10-fold cross-validation to ensure the models are generalizing well and stable over the data set rather than performing well on a single subsection of the train/test split.

Libraries used:

- tidymodels: used for model creation, cross-validation, and determining model performance
- tidyverse: used for plotting, data manipulation, and chaining operations
- probably: used for assessing threshold performance on the models
- discrim: used for the LDA/QDA models
- patchwork: used for combining plots
- doParallel: used for setting up parallel processing of the code to speed up performance.

Metrics utilized:

- specificity: $\frac{TN}{FP+TN}$

- sensitivity: $\frac{TP}{TP+FN}$

- accuracy: $\frac{TP+TN}{TP+TN+FP+FN}$

- precision: $\frac{TP}{TP+FP}$

Results

Logistic Regression Model

The logistic model was fit across the 10-fold cross-validation. The model seems to be generalizing well across the different folds with minimal variation across the training performance metrics as seen in *Figure 3*. The average performance metrics tell us that the initial model has very high sensitivity with little variation across the folds. Specificity is much lower with more variation in the results indicating is it not as stable. This can be interpreted as the model is performing well at classifying the blue tarps. However it is classifying many objects that are not blue tarps creating many false positives as seen in the lower specificity.

After performing threshold analysis, a threshold of 0.84 was selected to minimize the false positives while keeping a substantial portion of the true positives in the analysis. Seen in *Table 1*, Utilizing this threshold resulted in a precision rating of 0.996 meaning that 99.6% of the positive predictions in the model are true positives. The AUC score of 0.999 means that the model has near perfect classification abilities and is highly reliable in determining whether a data point is a blue tarp or not. Similar to the results from the training data, specificity is lower at 87.7% on the tesing data. The performance on the testing data in relation to the cross-validation of the training set indicates that the model is not overfitting on the training data

At the threshold of 0.84, 342 of the 390 blue tarps in the testing set are correctly identified by the model. Of the 12,259 not blue tarps, only 4 are false positives. Of all the blue tarps identified by the model, there is a 98.8% rate of aid workers time and resources not being wasted on futile searches. While 48 blue tarps were unidentified by the model (12.3% of total tarps), once the aid workers can provide resources to the 342 correctly identified refuges we can then expand the search. This will help us maximize a timely and precise search of refugees accounting for early and easy wins before expanding the search areas to get the last pockets of refugees.

LDA Model

The LDA Model has similar variations in the cross-validation performance as the logistic regression model seen in *Figure 3*. Specificity seems to be the only metric not performing at above 97%. The specificity across the 10-fold cross-validation of the training set is around 80% with slightly less variation within one standard error than the logistic regression.

The threshold selected was also 0.84 to keep the number of false positives at a minimum. At this threshold, the LDA model performs worse across all performance metrics and than logistic regression model, seen in *Table 1*. This seems to indicate that a linear decision boundary may not be the best fit for the data set being utilized.

At the selected threshold of 0.84, only 299 of the 390 blue tarps in the testing set are correctly identified by the model. Of the 12,259 not blue tarps, 88 were false positives. This would represent an increase of 2100% more false positives and could equate to numerous man hours and resources being wasted if we were to go with this model. 91 blue tarps were false negatives also resulting in more pockets of refuges being missed by this model.

QDA Model

The QDA model performs better than the LDA model for the 10-fold cross-validation but worse than the logistic regression. The main difference between all the models thus far has been the variation in specificity as sensitivity, precision, and accuracy have been comparable across the models. The QDA model also has

smaller variation within one standard error for specificity than the logistic regression but has a higher average specificity than the LDA model.

At the selected threshold of 0.85 on the testing data, the QDA model only has 2 false positives. However, the number of true positives predicted is 312 which results in 30 less pockets of refuges being rescued and the number of false negatives is 78 which is an increase of 62.5% more blue tarps being unidentified compared to the logistic regression model.

K-Nearest Neighbors Model

The KNN model performs better than the other models looking at the 10-fold cross validation. The sensitivity, accuracy, and precision are pretty similar to the other models, but the specificity performs better than LDA and QDA by a wide margin and performs better than logistic regression by 1%, which can make a pretty large impact in practice, since a higher specificity indicates more true positives, which means a higher chance of US military troops going to the correct location and saving valuable time when trying to find blue tarps with displaced persons.

The KNN model has 5 false positives and 43 false negatives, and finds more true positive blue tarp locations than all of the other models so far, narrowly beating out logistic regression with 5 more true positives, which means 5 more pockets of displaced persons found. 5 more true positives has significant real world significance in this case, since every additional correct pocket found is possibly many lives saved. So, finding more true blue tarp locations is the most important factor, since it means there will be more people saved and less errors made in the search, making the 20 neighbor KNN model a favorable option.

We tried several counts of neighbors for the KNN model and ended up choosing 20 neighbors due to it having the best performance metrics, importantly with the highest specificity since that is the metric with the most room for improvement. Lower than 20 neighbors (tested 10) indicated to possibly have signs of overfitting and higher K values had worse performance metrics on test data.

Penalized Logistic Regression Model

Table 1: Test Performance Metrics

Threshold	specificity	sensitivity	accuracy	precision	roc_auc	model
0.84	0.8769	0.9997	0.9959	0.9961	0.9988	logreg
0.84	0.7667	0.9928	0.9858	0.9926	0.9905	LDA
0.85	0.8000	0.9998	0.9937	0.9937	0.9973	QDA
0.85	0.8923	0.9996	0.9963	0.9966	0.9934	KNN
0.85	0.8692	0.9997	0.9957	0.9959	0.9989	Penalized LR

Metrics and Graphics

LDA Validation Results

QDA Validation Results

Logreg Cross Validation Result

KNN Validation Results

Penalized Logistic Regression Validation Results

Figure 3: Cross Validation Metrics

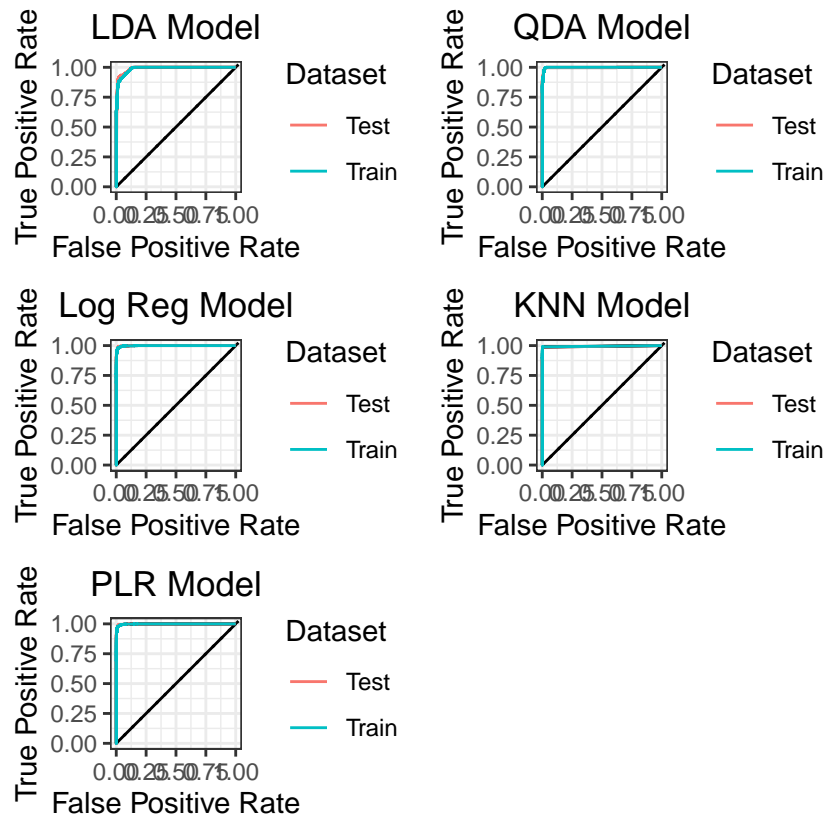Figure 4: ROC Curves

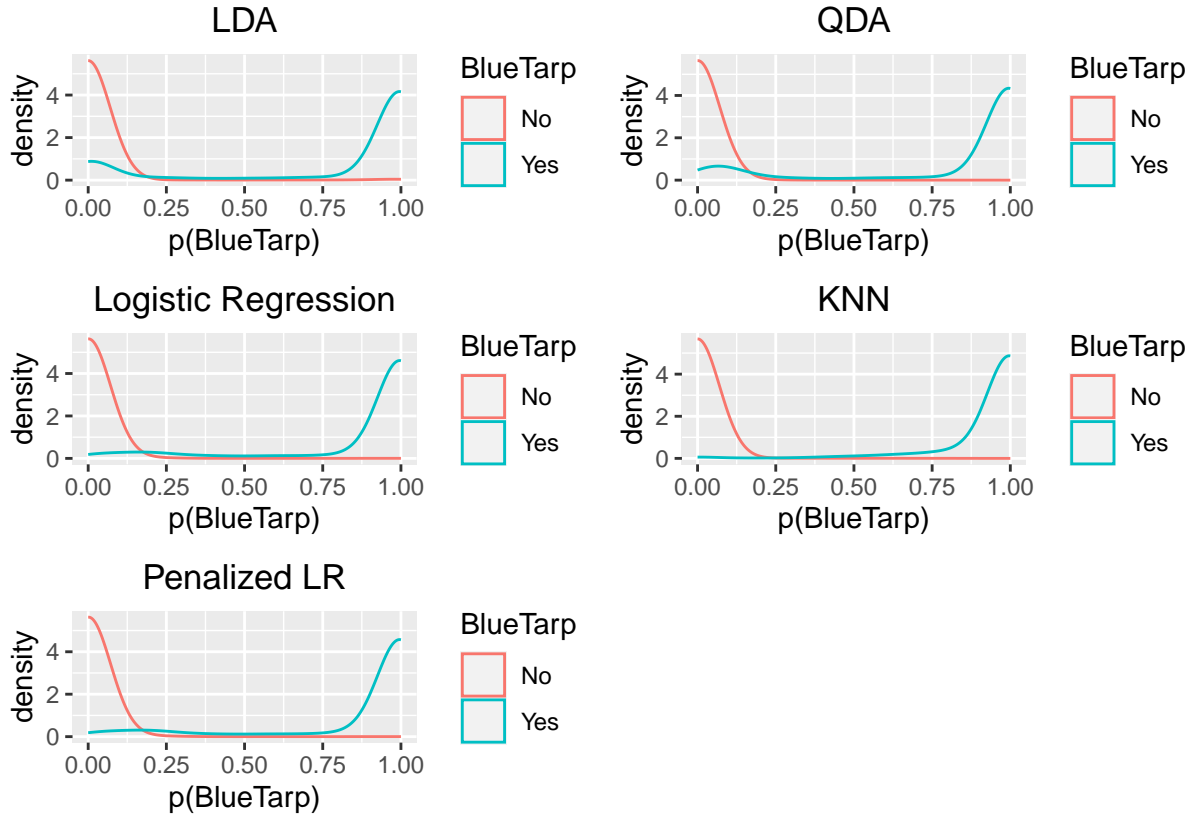Figure 5: Threshold Performance

```
##          Truth
## Prediction   No   Yes
##        No  12248   37
##       Yes    11  353

##          Truth
## Prediction   No   Yes
##        No  12257   78
##       Yes     2  312
```

Conclusion

Based on our analyses, we found that the Penalized Logistic Regression Model to perform the best out of the five models explored. Based on our performance metrics of specificity, sensitivity, accuracy, precision, and AUC, this model consistently scored either higher than, or similarly to, the other models. In particular, the PLR Model had the highest specificity, accuracy, and AUC out of all the models we tried, with values of 0.9051, 0.9962, and 0.9989, respectively. This model incorrectly classified 11 non-tarps as tarps, and incorrectly classified 37 tarps as non-tarps. We found this to be the best combination of false-positives and true-negatives in order to efficiently allocate resources in a rescue situation.

That being said, most of the models performed very well. There were no models which significantly outperformed, or underperformed, based on our performance metrics and qualitative analysis. We have high confidence in our results

Appendix

```r
knitr::opts_chunk$set(echo=FALSE)
knitr::opts_chunk$set(cache=TRUE, autodep=TRUE)
knitr::opts_chunk$set(fig.align="center", fig.pos="H")
# Set up Parallel Processing
library(doParallel)

cl <- makePSOCKcluster(parallel::detectCores(logical = FALSE))

registerDoParallel(cl)

# Load Libraries
library(tidyverse)
library(tidymodels)
library(probably)
library(discrim)
library(patchwork)

# Read in Data
haiti <- read_csv('https://gedeck.github.io/DS-6030/project/HaitiPixels.csv', show_col_types=FALSE) %>%
          mutate(BlueTarp= factor(ifelse(Class=="Blue Tarp", "Yes", "No"),labels=c("No", "Yes")))

# View Average Color of Each Class
haiti %>%
  group_by(Class) %>%
    summarize(R = mean(Red),
              G = mean(Green),
              B = mean(Blue))%>%
    mutate(hex = rgb(R, G, B, maxColorValue = 255))%>%
    ggplot(aes(x = 1, y = 1, fill = hex)) +
  geom_tile() +
  scale_fill_identity() +
  theme_void() +
  facet_wrap(~Class, ncol=5)
# Show different classifications in data set
haiti %>%
  group_by(Class) %>%
  summarize(count=n()) %>%
  mutate(percent_of_total = sprintf('%.1f%%', count / sum(count) * 100)) %>%
    ggplot(aes(x = reorder(Class, -count), y = count, fill=Class)) +
    geom_bar(stat = "identity") +
    scale_fill_manual(values=c('#A9BACD', '#C3B7A2' , '#F7E3B0','#B8A88C', '#4E4E3C')) +
    geom_text(aes(label = percent_of_total), vjust = -0.5, size = 3) +
    labs(x = "", y = "Instances") +
    theme_minimal()+
    guides(fill = "none")

# Set seed
set.seed(81718)

# Create initial split for 80/20 with stratified sampling on BlueTarp
haiti_split <- initial_split(haiti, prop=.8, strat=BlueTarp)
```

```r
# Create training data set
train_data <- training(haiti_split)

# Create test data set
test_data <- testing(haiti_split)

# Set up 10-fold cross-validation
resamples <- vfold_cv(train_data, v=10, strata=BlueTarp)

# Set settings for control resamples
cv_control <- control_resamples(save_pred=TRUE)

# Define performance metrics
performance_metrics <- metric_set(specificity, sensitivity, accuracy, precision)


get_ROC_plot <- function(model, train_data, test_data, model_name){
  # Augment train and test data with predicted probabilities
  roc_train <- augment(model, train_data) %>%
    roc_curve(truth = BlueTarp, .pred_Yes, event_level = "second") %>%
    mutate(Dataset = "Train")

  roc_test <- augment(model, test_data) %>%
    roc_curve(truth = BlueTarp, .pred_Yes, event_level = "second") %>%
    mutate(Dataset = "Test")

  # Combine train and test ROC curve data
  roc_data <- bind_rows(roc_train, roc_test)

    # Plot ROC curves for train and test data with different colors
  autoplot(roc_data) +
    geom_line(aes(x = 1 - specificity, y = sensitivity, color = Dataset))+
    labs(title = model_name, x = "False Positive Rate", y = "True Positive Rate")+
    theme(plot.title = element_text(hjust = 0.5))
}

# Create Function to Visual Train Metrics
visualize_training <- function(fit_resample_results, title){

    aggregate_metrics <-  bind_rows(fit_resample_results$.metrics) %>%
        group_by(.metric) %>%
        summarize(Mean = mean(.estimate),
                  std_err = sd(.estimate) / sqrt(n()))%>%
        rename(Metric=.metric)

    aggregate_metrics %>%
      ggplot(aes(x=Mean, y=Metric, xmin=Mean-std_err, xmax=Mean+std_err)) +
      geom_point() +
      geom_linerange() +
      ggtitle(title) +
      theme(plot.title = element_text(hjust = 0.5))
}
```

```r
# Create Function to visualize distributions
distribution_graph <- function(model, data, model_name) {
                        model %>%
                            augment(data) %>%
                                ggplot(aes(x=.pred_Yes, color=BlueTarp)) +
                                geom_density(bw=0.07) +
                                labs(x='p(BlueTarp)', title=model_name) +
                                theme(plot.title = element_text(hjust = 0.5))
}

# Test Thresholds
performance_func_1 <- function(model, data){
                    threshold_perf(model %>% augment(train_data),
                                    BlueTarp,
                                    .pred_Yes,
                                    thresholds = seq(0.01, 0.85, 0.01), event_level="second",
                                    metrics=performance_metrics)
}

# Pick best precision as Threshold
max_precision <- function(performance_data){
            performance_data %>%
                filter(.metric == 'precision') %>%
                filter(.estimate == max(.estimate))
}


# Create Formula
formula <- BlueTarp ~ Red + Green + Blue

# Create Recipe
rec <- recipe(formula, data=train_data) %>%
    step_normalize(all_numeric_predictors())

# Create Log Model
logreg_model <- logistic_reg() %>%
    set_engine("glm") %>%
    set_mode("classification")

# define and execute the cross-validation workflow
logreg_wf <- workflow() %>%
    add_model(logreg_model) %>%
    add_recipe(rec)

# Cross Validate Model
logreg_fit_cv <- logreg_wf %>%
                    fit_resamples(resamples=resamples, control=cv_control, metrics=performance_metrics)

# Visualize logreg Fit
logreg_cv_viz <- visualize_training(logreg_fit_cv, "Logreg Cross Validation Results")

# Fit Model
logreg_model_fit <- logreg_wf %>% fit(train_data)
```

```r
# Get Performance Thresholds
logreg_threshold_performance <- performance_func_1(logreg_model_fit)

# Run Model on Test Data
logreg_results <- logreg_model_fit %>% augment(test_data)

# Change Pred Class metric based on threshold testing
logreg_results$.threshold_pred_class <- as.factor(ifelse(logreg_results$.pred_Yes >= max_precision(logre

# View results before and after threshold picking
performance_table <-  performance_metrics(logreg_results, truth=BlueTarp, estimate=.threshold_pred_clas
                            bind_rows(roc_auc(logreg_results, truth=BlueTarp, .pred_Yes, event_lev
                            mutate(Threshold = max_precision(logreg_threshold_performance)$.thresl
                              dplyr::select(c(Threshold, .metric, .estimate)) %>%
                                pivot_wider(names_from = .metric, values_from = .estimate, id_col
                                mutate(model="logreg")


# Create LDA Model
lda_model <- discrim_linear(mode="classification") %>%
                set_engine("MASS")

# Create Workflow
lda_wf <- workflow()%>%
            add_model(lda_model)%>%
            add_recipe(rec)

# Create Validation Metric Set
lda_wf_fit_cv <- lda_wf %>%
                    fit_resamples(resamples=resamples, control=cv_control, metrics=performance_metrics)

# Visualize CV Results Fit
lda_cv_viz <- visualize_training(lda_wf_fit_cv, "LDA Validation Results")

# Fit LDA Model
lda_model_fit <- lda_wf %>% fit(train_data)

# Run Model on Test Data
lda_results <- lda_model_fit %>% augment(test_data)

# Get Performance Thresholds
lda_threshold_performance <- performance_func_1(lda_model_fit)

# Change Pred Class metric based on threshold testing
lda_results$.threshold_pred_class <- as.factor(ifelse(lda_results$.pred_Yes >= max_precision(lda_thresho

# View results before and after threshold picking
performance_table <-  bind_rows(performance_table,
                            performance_metrics(lda_results, truth=BlueTarp, estimate=.threshold_pre
                            bind_rows(roc_auc(lda_results, truth=BlueTarp, .pred_Yes, event_level
                            mutate(Threshold = max_precision(lda_threshold_performance)$.threshold
                            dplyr::select(c(Threshold, .metric, .estimate)) %>%
                                pivot_wider(names_from = .metric, values_from = .estimate, id_col
```

```r
                                    mutate(model="LDA")
                    )

# Create QDA Model
qda_model <- discrim_quad(mode="classification") %>%
                set_engine("MASS")

# Create Workflow
qda_wf <- workflow()%>%
            add_model(qda_model)%>%
            add_recipe(rec)

# Create Validation Metric Set
qda_wf_fit_cv <- qda_wf %>%
                    fit_resamples(resamples=resamples, control=cv_control, metrics=performance_metrics)

# Visualize QDA Fit
qda_cv_viz <- visualize_training(qda_wf_fit_cv, "QDA Validation Results")

# Fit QDA Model
qda_model_fit <- qda_wf %>% fit(train_data)

# Run Model on Test Data
qda_results <- qda_model_fit %>% augment(test_data)

# Get Performance Thresholds
qda_threshold_performance <- performance_func_1(qda_model_fit)

# Change Pred Class metric based on threshold testing
qda_results$.threshold_pred_class <- as.factor(ifelse(qda_results$.pred_Yes >= max_precision(qda_thresho


# View results before and after threshold picking
performance_table <-  bind_rows(performance_table,
                            performance_metrics(qda_results, truth=BlueTarp, estimate=.threshold_pr
                                bind_rows(roc_auc(qda_results, truth=BlueTarp, .pred_Yes, event_level
                                mutate(Threshold = max_precision(qda_threshold_performance)$.threshol
                                dplyr::select(c(Threshold, .metric, .estimate)) %>%
                                    pivot_wider(names_from = .metric, values_from = .estimate, id_col
                                    mutate(model="QDA")
                    )

# performance_table %>%
#       knitr::kable(digits=4, caption='Test Performance Metrics')

# Create KNN Model
knn_model <- nearest_neighbor(neighbors=tune()) %>%
    set_mode("classification") %>%
    set_engine("kknn")


# Create Workflow
knn_wf <- workflow() %>%
```

```r
    add_model(knn_model) %>%
    add_recipe(rec)

# Set neighbor range
knn_model_params <- extract_parameter_set_dials(knn_wf) %>%
  update(neighbors=neighbors(range=c(5, 50)))

# Tune neighbors param
knn_tune_results <- tune_grid(knn_wf, resamples=resamples,
                              control=cv_control,
                              grid=grid_random(knn_model_params, size=22))

# show_best(knn_tune_results, metric="accuracy", n=1)
# show_best(knn_tune_results, metric="roc_auc", n=1)

# Tune the model based on the best grid search result
knn_tuned_model <- knn_wf %>% finalize_workflow(select_best(knn_tune_results, metric="roc_auc"))

# Create Validation Metric Set
knn_wf_fit_cv <- knn_tuned_model %>%
    fit_resamples(resamples, control=control_resamples(save_pred=TRUE),metrics=performance_metrics)


# Visualize KNN Fit
knn_cv_viz <- visualize_training(knn_wf_fit_cv, "KNN Validation Results")


# Fit KNN Model
knn_model_final_fit <- knn_tuned_model %>% fit(train_data)

# Run Model on Test Data
knn_results <- knn_model_final_fit %>% augment(test_data)


# Get Performance Thresholds
knn_threshold_performance <- performance_func_1(knn_model_final_fit)

# Change Pred Class metric based on threshold testing
knn_results$.threshold_pred_class <- as.factor(ifelse(knn_results$.pred_Yes >= max_precision(knn_thresho

# View results before and after threshold picking
performance_table <-  bind_rows(performance_table,
                                performance_metrics(knn_results, truth=BlueTarp, estimate=.threshold_pre
                                  bind_rows(roc_auc(knn_results, truth=BlueTarp, .pred_Yes, event_level
                                  mutate(Threshold = max_precision(knn_threshold_performance)$.threshold
                                  dplyr::select(c(Threshold,
                                                 .metric, .estimate)) %>%
                                    pivot_wider(names_from = .metric, values_from = .estimate, id_cols
                                    mutate(model="KNN")
                                )
#
# performance_table %>%
```

```r
#       knitr::kable(digits=4, caption='Test Performance Metrics')

# Create Penalized Log Model
elasticnet_spec <- logistic_reg(engine="glmnet", mode="classification",
                                penalty=tune(), mixture=tune())

# Create Workflow
elasticnet_wf <- workflow() %>%
    add_model(elasticnet_spec) %>%
    add_recipe(rec)

# Testing a range of parameter values
parameters <- extract_parameter_set_dials(elasticnet_wf) %>%
    update(
        penalty=penalty(c(-20, -10)),
        mixture=mixture(c(0, 10))
    )


# Identify best model based on searching the parameter space
tune_results <- tune_grid(elasticnet_wf,
                          resamples=resamples,
                          grid=grid_random(parameters, size=200),
                          metrics=performance_metrics)

# show_best(tune_results, metric='accuracy', n=1) %>%
#     knitr::kable()
#
# autoplot(tune_results)

# Finalize workflow with best model and fit model
elasticnet_model_tuned <- elasticnet_wf %>%
    finalize_workflow(select_best(tune_results, metric="accuracy")) %>%
  fit(train_data)

# Get preds with test data
elastic_results <- elasticnet_model_tuned %>% augment(test_data)



# Visualize Penalized Log Fit
elastic_cv_viz <- visualize_training(tune_results,
                                "Penalized Logistic
                                Regression Validation Results")

# Run Model on Test Data
# elastic_model_fit <- elasticnet_wf %>% fit(train_data)



# Create Validation Metric Set
# elasticnet_wf_fit_cv <- elasticnet_model_tuned %>%
#     fit_resamples(resamples, control=control_resamples(save_pred=TRUE),metrics=performance_metrics)
```

```r
# Fit EN Model
# elasticnet_model_final_fit <- elasticnet_model_tuned %>% fit(train_data)

# Get threshold performance
en_threshold_performance <- performance_func_1(elasticnet_model_tuned)

# Change Pred Class metric based on threshold testing
elastic_results$.threshold_pred_class <- as.factor(ifelse(elastic_results$.pred_Yes >= max_precision(en_

# tune_results %>%
#     show_best(metric='accuracy') %>%
#     select(-.config) %>%
#     knitr::kable()
#
# test_predictions <- elasticnet_model_tuned %>%
#   predict(test_data) %>%
#   bind_cols(test_data)
#
# test_metrics <- test_predictions %>%
#   performance_metrics(truth = BlueTarp, estimate = .pred_class)

# View results before and after threshold picking
performance_table <-  bind_rows(performance_table,
                        performance_metrics(elastic_results, truth=BlueTarp, estimate=.threshold
                          bind_rows(roc_auc(elastic_results, truth=BlueTarp, .pred_Yes, event_le
                          mutate(Threshold = max_precision(en_threshold_performance)$.threshold
                          dplyr::select(c(Threshold,
                                          .metric, .estimate)) %>%
                              pivot_wider(names_from = .metric, values_from = .estimate, id_cols
                              mutate(model="Penalized LR")
                        )

performance_table %>%
      knitr::kable(digits=4, caption='Test Performance Metrics')

# Visualize Cross validation Metrics
lda_cv_viz + qda_cv_viz + logreg_cv_viz + knn_cv_viz + elastic_cv_viz + plot_layout(ncol=2)
log_roc <- get_ROC_plot(logreg_model_fit, train_data, test_data, "Log Reg Model")
lda_roc <- get_ROC_plot(lda_model_fit, train_data, test_data, "LDA Model")
qda_roc <- get_ROC_plot(qda_model_fit, train_data, test_data, "QDA Model")
knn_roc <- get_ROC_plot(knn_model_final_fit, train_data, test_data, "KNN Model")
elastic_roc <- get_ROC_plot(elasticnet_model_tuned, train_data, test_data, "PLR Model")

# Visualize ROC plots
lda_roc + qda_roc + log_roc + knn_roc + elastic_roc +  plot_layout(ncol=2)


log_reg_distribution <- distribution_graph(logreg_model_fit, train_data, "Logistic Regression")
lda_distribution <- distribution_graph(lda_model_fit, train_data, "LDA")
qda_distribution <- distribution_graph(qda_model_fit, train_data, "QDA")
knn_distribution <- distribution_graph(knn_model_final_fit, train_data, "KNN")
en_distribution <- distribution_graph(elasticnet_model_tuned , train_data, "Penalized LR")
```

```r
# Visualize Distributions
lda_distribution + qda_distribution + log_reg_distribution + knn_distribution + en_distribution + plot_l

# Create Confusion Matrixes to reference TP, FP,TN, FP in discussion of results
logreg_conf_matrix <- conf_mat(logreg_results, estimate=.threshold_pred_class, truth=BlueTarp)
lda_conf_matrix <- conf_mat(lda_results, estimate=.threshold_pred_class, truth=BlueTarp)
qda_conf_matrix <- conf_mat(qda_results, estimate=.threshold_pred_class, truth=BlueTarp)
knn_conf_matrix <- conf_mat(knn_results, estimate=.threshold_pred_class, truth=BlueTarp)
en_conf_matrix <- conf_mat(elastic_results, estimate=.pred_class, truth=BlueTarp)


en_conf_matrix
qda_conf_matrix
```