

Disaster-Relief-Project-Part-1

Mahin Ganesan, Wyatt Priddy, and Taylor Tucker

2024-03-13

Introduction

In the aftermath of a devastating earthquake that rattled Haiti, displaced people are living in makeshift shelters awaiting support from aid workers. The aid workers, mainly from the United States military, are trying to reach the dispersed camps. With communication lines being down and the terrain being impassable, there are challenges in providing relief in a time- sensitive manner.

It is known that the makeshift shelters are largely constructed with blue tarps, therefore the Rochester Institute of Technology deployed airplanes to collect high resolution geo-referenced imagery. This will help us identify where displaced persons are based on the tarps.

To determine where to allocate aid, we need to use data mining against the thousands of photos taken each day which human eyes can not efficiently filter through. Determining the location in a timely manner will be paramount to rendering aid successfully and saving human life.

The primary aim of this experiment is to evaluate the efficacy of various classification algorithms in accurately and promptly identifying the presence of makeshift shelters and, by extension, the displaced persons residing within them. By harnessing the power of machine learning and image analysis, our objective is to develop a robust algorithm capable of rapidly scanning through the imagery data, pinpointing areas of interest, and facilitating timely intervention by rescue teams.

To facilitate efficient rescue efforts, we will need to determine a threshold where the number of false positives is minimal. This may not necessarily be the model that has the highest performance in accuracy but rather the highest performance with precision, or the proportion of true positives that are correctly identified by the model out of all true positives and false positive. If there are additional aid resources after rescuing all persons identified from our models, we can loosen our model specifications to classify more objects as blue tarps in an attempt to expand our search efforts.

Data

Team members have assisted our mission by classifying training data consisting of 63,241 data points for our investigation. There are 5 classifications that have been assigned to the pixel level data:

1. Blue Tarp
2. Rooftop
3. Soil
4. Various Non-Tarp
5. Vegetation

As expected when trying to find a needle in a haystack, our representation of misplaced persons (blue tarps) makes up only a small portion of the data set. Just 3.2% of the total data set, or 2022 records, are classified

as **blue tarp**. After inspecting the average color of the classes, it becomes apparent that even though blue tarps are a small section of the data their distinction in color should set them apart.




Figure 1: Average Color of Class

Vegetation and soil cover over 73% of the pictures as to be expected of pictures that largely will include countryside.




Figure 2: Distribution of Classifications in Training Set

Description of Methodology

Given that our sample data set only has 3.2% of the target class, we will use stratification to account for the imbalance in our test and train split. We are looking to ensure that there is even representation between our two sets to ensure a robust analysis when training our models and subsequently predicting on the test data. We are using an 80/20 split on the training versus testing data. Since the data set is relatively large and we are using stratification to ensure an even population of blue tarps across the split, leaving 20% of the data for hold-out is appropriate.

Each model is being trained with 10-fold cross-validation to ensure the models are generalizing well and stable over the data set rather than performing well on a single subsection of the train/test split.

Libraries used:

- `tidymodels`: used for model creation, cross-validation, and determining model performance
- `tidyverse`: used for plotting, data manipulation, and chaining operations

- probably: used for assessing threshold performance on the models
- doParallel: used for setting up parallel processing of the code to speed up performance.

Metrics utilized:

- specificity: $\frac{TN}{FP+TN}$
- sensitivity: $\frac{TP}{TP+FN}$
- accuracy: $\frac{TP+TN}{TP+TN+FP+FN}$
- precision: $\frac{TP}{TP+FP}$

Results

Logistic Regression Model

The logistic model was fit across the 10-fold cross-validation. The model seems to be generalizing well across the different folds with minimal variation across the training performance metrics. The average performance metrics tell us that the initial model has very high sensitivity with little variation across the folds. Specificity is much lower with more variation in the results indicating it is not as stable. This can be interpreted as the model is performing well at classifying the blue tarps. However it is classifying many objects that are not blue tarps creating many false positives as seen in the lower specificity.




Figure 3: Logistic Regression CV Results


After performing threshold analysis, a threshold of 0.84 was selected to minimize the false positives while keeping a substantial portion of the true positives in the analysis. Utilizing this threshold resulted in a precision rating of 0.996 meaning that 99.6% of the positive predictions in the model are true positives. The AUC score of 0.999 means that the model has near perfect classification abilities and is highly reliable in determining whether a data point is a blue tarp or not. Similar to the results from the training data, specificity is lower at 87.7% on the testing data. The performance on the testing data in relation to the cross-validation of the training set indicates that the model is not overfitting on the training data

Table 1: Test Performance Metrics

Threshold	specificity	sensitivity	accuracy	precision	roc_auc
0.84	0.8769	0.9997	0.9959	0.9961	0.9988

At the threshold of 0.84, 342 of the 390 blue tarps in the testing set are correctly identified by the model. Of the 12259 not blue tarps, only 4 are false positives. Of all the blue tarps identified by the model, there is a 98.8% rate of aid workers time and resources not being wasted on futile searches. While 48 blue tarps were unidentified by the model (12.3%), once the aid workers can provide resources to the 342 correctly identified refugees we can then expand the search. This will help us maximize a timely and precise search of refugees accounting for early and easy wins before expanding the search areas to get the last pockets of refugees.

TODO PUT THIS IN GRAPH WITH OTHER MODELS ONCE THEY ARE DEVELOPED USING PATCHWORK



Conclusion

TEXT

Appendix

```
knitr::opts_chunk$set(echo=FALSE)
knitr::opts_chunk$set(cache=TRUE, autodep=TRUE)
knitr::opts_chunk$set(fig.align="center", fig.pos="H")
# Set up Parallel Processing
library(doParallel)

cl <- makePSOCKcluster(parallel::detectCores(logical = FALSE))

registerDoParallel(cl)

# Load Libraries
library(tidyverse)
library(tidymodels)
library(probably)

# Read in Data
haiti <- read_csv('https://gedeck.github.io/DS-6030/project/HaitiPixels.csv', show_col_types=FALSE) %>%
  mutate(BlueTarp= factor(ifelse(Class=="Blue Tarp", "Yes", "No"), labels=c("No", "Yes")))

# View Average Color of Each Class
haiti %>%
  group_by(Class) %>%
  summarize(R = mean(Red),
            G = mean(Green),
            B = mean(Blue))%>%
  mutate(hex = rgb(R, G, B, maxColorValue = 255))%>%
  ggplot(aes(x = 1, y = 1, fill = hex)) +
  geom_tile() +
  scale_fill_identity() +
  theme_void() +
  facet_wrap(~Class, ncol=5)

# Show different classifications in data set
haiti %>%
  group_by(Class) %>%
  summarize(count=n()) %>%
  mutate(percent_of_total = sprintf('%.1f%%', count / sum(count) * 100)) %>%
  ggplot(aes(x = reorder(Class, -count), y = count, fill=Class)) +
  geom_bar(stat = "identity") +
  scale_fill_manual(values=c('#A9BACD', '#C3B7A2', '#F7E3B0', '#B8A88C', '#4E4E3C')) +
  geom_text(aes(label = percent_of_total), vjust = -0.5, size = 3) +
  labs(x = "", y = "Instances") +
  theme_minimal()+
  guides(fill = "none")

# Set seed
set.seed(81718)

# Create initial split for 80/20 with stratified sampling on BlueTarp
haiti_split <- initial_split(haiti, prop=.8, strat=BlueTarp)
```

```

# Create training data set
train_data <- training(haiti_split)

# Create test data set
test_data <- testing(haiti_split)

# Set up 10-fold cross-validation
resamples <- vfold_cv(train_data, v=10, strata=BlueTarp)

# Set settings for control resamples
cv_control <- control_resamples(save_pred=TRUE)

# Define performance metrics
performance_metrics <- metric_set(specificity, sensitivity, accuracy, precision)

get_ROC_plot <- function(model, train_data, test_data, model_name){
  # Augment train and test data with predicted probabilities
  roc_train <- augment(model, train_data) %>%
    roc_curve(truth = BlueTarp, .pred_Yes, event_level = "second") %>%
    mutate(Dataset = "Train")

  roc_test <- augment(model, test_data) %>%
    roc_curve(truth = BlueTarp, .pred_Yes, event_level = "second") %>%
    mutate(Dataset = "Test")

  # Combine train and test ROC curve data
  roc_data <- bind_rows(roc_train, roc_test)

  # Plot ROC curves for train and test data with different colors
  autoplot(roc_data) +
    geom_line(aes(x = 1 - specificity, y = sensitivity, color = Dataset)) +
    labs(title = model_name, x = "False Positive Rate", y = "True Positive Rate") +
    theme(plot.title = element_text(hjust = 0.5))
}

# Create Function to Visual Train Metrics
visualize_training <- function(fit_resample_results, title){

  aggregate_metrics <- bind_rows(logreg_fit_cv$.metrics) %>%
    group_by(.metric) %>%
    summarize(Mean = mean(.estimate),
              std_err = sd(.estimate) / sqrt(n())) %>%
    rename(Metric=.metric)

  aggregate_metrics %>%
    ggplot(aes(x=Mean, y=Metric, xmin=Mean-std_err, xmax=Mean+std_err)) +
    geom_point() +
    geom_linerange() +
    ggtitle(title) +
    theme(plot.title = element_text(hjust = 0.5))
}

```

```

# Test Thresholds
performance_func_1 <- function(model, data){
  threshold_perf(model %>% augment(train_data),
    BlueTarp,
    .pred_Yes,
    thresholds = seq(0.01, 0.85, 0.01), event_level="second",
    metrics=performance_metrics)
}

# Pick best precision as Threshold
max_precision <- function(performance_data){
  performance_data %>%
    filter(.metric == 'precision') %>%
    filter(.estimate == max(.estimate))
}

# Create Formula
formula <- BlueTarp ~ Red + Green + Blue

# Create Recipe
rec <- recipe(formula, data=train_data) %>%
  step_normalize(all_numeric_predictors())

# Create Log Model
logreg_model <- logistic_reg() %>%
  set_engine("glm") %>%
  set_mode("classification")

# define and execute the cross-validation workflow
logreg_wf <- workflow() %>%
  add_model(logreg_model) %>%
  add_recipe(rec)

# Cross Validate Model
logreg_fit_cv <- logreg_wf %>%
  fit_resamples(resamples=resamples, control=cv_control, metrics=performance_metrics)

# Visualize logreg Fit
visualize_training(logreg_fit_cv, "Logreg Cross Validation Results")

# Fit Model
logreg_model <- logreg_wf %>% fit(train_data)

# Get Performance Thresholds
threshold_performance <- performance_func_1(logreg_model)

# Run Model on Test Data
logreg_results <- logreg_model %>% augment(test_data)

# Change Pred Class metric based on threshold testing
logreg_results$.threshold_pred_class <- as.factor(ifelse(logreg_results$.pred_Yes >= max_precision(thre

```

```

# View results before and after threshold picking
performance_table <- performance_metrics(logreg_results, truth=BlueTarp, estimate=.threshold_pred_class,
                                         bind_rows(roc_auc(logreg_results, truth=BlueTarp, .pred_Yes, event_level="Yes",
                                                             estimate=.threshold_pred_class, event_level="Yes",
                                                             event_level="No", event_level="No") %>%
                                         mutate(Threshold = max_precision(threshold_performance)$threshold) %>%
                                         select(c(Threshold, .metric, .estimate)) %>%
                                         pivot_wider(names_from = .metric, values_from = .estimate, id_col = Threshold))

performance_table %>%
  knitr::kable(digits=4, caption='Test Performance Metrics')
logreg_conf_matrix <- conf_mat(logreg_results, estimate=.threshold_pred_class, truth=BlueTarp)
get_ROC_plot(logreg_model, train_data, test_data, "Log Reg Model")

```